

```
In [1]: %matplotlib inline
      ## import libraries
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt
      from matplotlib import style
      import seaborn as sns
```

```
In [2]: data = pd.read_csv('health care diabetes.csv')
```

```
In [3]: data.head()
```

Out[3]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [4]: data.isnull().any()
```

Out[4]:

Pregnancies	False
Glucose	False
BloodPressure	False
SkinThickness	False
Insulin	False
BMI	False
DiabetesPedigreeFunction	False
Age	False
Outcome	False

dtype: bool

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [6]: Positive = data[data['Outcome']==1]
      Positive.head(5)
```

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
2	8	183	64	0	0	23.3	0.672	32	1
4	0	137	40	35	168	43.1	2.288	33	1
6	3	78	50	32	88	31.0	0.248	26	1
8	2	197	70	45	543	30.5	0.158	53	1

```
In [7]: data['Glucose'].value_counts().head(7)
```

Out[7]:


Glucose	
99	17
100	17
111	14
129	14
125	14
106	14
112	13

Name: count, dtype: int64

```
In [8]: plt.hist(data['Glucose'])
```

```
In [8]: plt.hist(data['Glucose'])
```

```
Out[8]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
array([ 0., 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
179.1, 199. ]),
<BarContainer object of 10 artists>)
```




```
In [9]: data['BloodPressure'].value_counts().head(7)
```

```
Out[9]: BloodPressure
70      57
74      52
78      45
68      45
72      44
64      43
80      40
Name: count, dtype: int64
```

```
In [10]: plt.hist(data['BloodPressure'])
```

```
Out[10]: (array([ 35.,  1.,  2., 13., 107., 261., 243., 87., 14.,  5.]),
array([ 0., 12.2, 24.4, 36.6, 48.8, 61., 73.2, 85.4, 97.6,
109.8, 122. ]),
<BarContainer object of 10 artists>)
```




```
In [11]: data['SkinThickness'].value_counts().head(7)
```

```
Out[11]: SkinThickness
0      227
32      31
30      27
27      23
23      22
33      20
28      20
Name: count, dtype: int64
```

```
In [12]: plt.hist(data['SkinThickness'])
```

```
Out[12]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),
array([ 0., 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
<BarContainer object of 10 artists>)
```




```
In [13]: data['Insulin'].value_counts().head(7)
```

```
Out[13]: Insulin
0      374
105     11
130      9
140      9
120      8
94       7
180      7
Name: count, dtype: int64
```

```
In [14]: plt.hist(data['Insulin'])
```

```
Out[14]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
array([ 0., 84.6, 169.2, 253.8, 338.4, 423., 507.6, 592.2, 676.8,
       761.4, 846. ]),
<BarContainer object of 10 artists>)
```




```
In [15]: data['BMI'].value_counts().head(7)
```

```
Out[15]: BMI
32.0      13
31.6      12
31.2      12
0.0       11
32.4      10
33.3      10
30.1       9
Name: count, dtype: int64
```

```
In [16]: plt.hist(data['BMI'])
```

```
Out[16]: (array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),
array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
        60.39, 67.1 ]),
<BarContainer object of 10 artists>)
```



```
In [17]: data.describe().transpose()
```

```
Out[17]:
```


	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	3.0000	6.00000	17.00
Glucose	768.0	120.894531	31.972618	0.000	99.00000	117.0000	140.25000	199.00
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	72.0000	80.00000	122.00
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	23.0000	32.00000	99.00
Insulin	768.0	79.799479	115.244002	0.000	0.00000	30.5000	127.25000	846.00
BMI	768.0	31.992578	7.884160	0.000	27.30000	32.0000	36.60000	67.10
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	0.3725	0.62625	2.42
Age	768.0	33.240885	11.760232	21.000	24.00000	29.0000	41.00000	81.00
Outcome	768.0	0.348958	0.476951	0.000	0.00000	0.0000	1.00000	1.00

```
In [ ]:
```

Week 2

```
In [18]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
Out[18]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
        14.,  7.,  3.,  1.,  1.,  0.,  1.]),
array([ 0. ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
        26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
        53.68 , 57.035, 60.39 , 63.745, 67.1 ]),
[<matplotlib.patches.Polygon at 0x1f3cd35d650>])
```




```
In [19]: Positive['BMI'].value_counts().head(7)
```

```
Out[19]: BMI
32.9    8
31.6    7
33.3    6
31.2    5
30.5    5
32.0    5
34.3    4
Name: count, dtype: int64
```

```
In [20]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
Out[20]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
 array([ 0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199. ]),
 [<matplotlib.patches.Polygon at 0x1f3cf58ac10>])
```




```
In [21]: Positive['Glucose'].value_counts().head(7)
```

```
Out[21]: Glucose
125    7
128    6
129    6
115    6
158    6
146    5
124    5
Name: count, dtype: int64
```

```
In [22]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
Out[22]: (array([16., 0., 0., 0., 0., 1., 0., 1., 6., 6., 19., 37., 56.,
        36., 41., 31., 7., 4., 4., 3.]),
array([ 0.,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
        51.3, 57., 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
        102.6, 108.3, 114. ]),
[<matplotlib.patches.Polygon at 0x1f3cf6bac10>])
```




```
In [23]: Positive['BloodPressure'].value_counts().head(7)
```

```
Out[23]: BloodPressure
70    23
76    18
78    17
74    17
72    16
0     16
80    13
Name: count, dtype: int64
```

```
In [24]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
Out[24]: (array([88., 1., 4., 10., 18., 30., 41., 34., 23., 15., 1., 1., 1.,
        0., 0., 0., 0., 0., 0., 1.]),
array([ 0.,  4.95, 9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
        44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
        89.1 , 94.05, 99. ]),
[<matplotlib.patches.Polygon at 0x1f3cf89ef50>])
```




```
In [25]: Positive['SkinThickness'].value_counts().head(7)
```

```
Out[25]: SkinThickness
0      88
32     14
30      9
33      9
39      8
37      8
36      8
Name: count, dtype: int64
```

```
In [26]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
Out[26]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
        5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
 array([ 0.,  42.3,  84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
        380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
        761.4, 803.7, 846. ]),
 [<matplotlib.patches.Polygon at 0x1f3cf91f9d0>])
```




```
In [27]: Positive['Insulin'].value_counts().head(7)
```

```
Out[27]: Insulin
0      138
130     6
180     4
175     3
156     3
185     2
194     2
Name: count, dtype: int64
```


```
In [28]: #Scatter plot
```

```
In [29]: BloodPressure = Positive['BloodPressure']
Glucose = Positive['Glucose']
SkinThickness = Positive['SkinThickness']
Insulin = Positive['Insulin']
BMI = Positive['BMI']
```


```
In [30]: plt.scatter(BloodPressure, Glucose, color='b')
plt.xlabel('BloodPressure')
plt.ylabel('Glucose')
plt.title('BloodPressure & Glucose')
plt.show()
```


```
In [31]: g = sns.scatterplot(x= "Glucose" ,y= "BloodPressure",  
                             hue="Outcome",  
                             data=data);
```



```
In [32]: B = sns.scatterplot(x= "BMI" ,y= "Insulin",  
                             hue="Outcome",  
                             data=data);
```



```
In [33]: S=sns.scatterplot(x= "SkinThickness" ,y= "Insulin",
hue="Outcome",
data=data);
```




```
In [34]: ### correlation matrix
data.corr()
```

Out[34]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0


```
In [35]: ### create correlation heat map
sns.heatmap(data.corr())
```

Out[35]: <Axes: >




```
In [36]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True,cmap='viridis') ### gives correlation value
```

```
Out[36]: <Axes: >
```



```
In [37]: plt.subplots(figsize=(8,8))
sns.heatmap(data.corr(),annot=True) ### gives correlation value
```

```
Out[37]: <Axes: >
```



```
In [38]: # Logistic Regrestion and model building
```

```
In [39]: data.head(5)
```

```
Out[39]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [40]: features = data.iloc[:,[0,1,2,3,4,5,6,7]].values
label = data.iloc[:,8].values
```

```
In [41]: #Train test split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(features,
                                                label,
                                                test_size=0.2,
                                                random_state =10)
```

```
In [42]: #Create model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,y_train)
```

```
Out[42]:
```

▼ LogisticRegression
 LogisticRegression()

```
In [43]: print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7719869706840391
0.7662337662337663
```

```
In [44]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(label,model.predict(features))
cm
```

```
Out[44]: array([[446,  54],
               [122, 146]], dtype=int64)
```

```
In [45]: from sklearn.metrics import classification_report
print(classification_report(label,model.predict(features)))
```

```

              precision    recall  f1-score   support

     0       0.79         0.89         0.84         500
     1       0.73         0.54         0.62         268


 accuracy          0.77         0.77         0.77         768
 macro avg         0.76         0.72         0.73         768
 weighted avg         0.77         0.77         0.76         768
```

```
In [46]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
```

```
AUC: 0.837
```

```
Out[46]: [ <matplotlib.lines.Line2D at 0x1f3d21d97d0>]
```



```
In [47]: #Applying Decision Tree Classifier
from sklearn.tree import DecisionTreeClassifier
model3 = DecisionTreeClassifier(max_depth=5)
model3.fit(X_train,y_train)
```

```
Out[47]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)
```

```
In [48]: model3.score(X_train,y_train)
```

```
Out[48]: 0.8289902280130294
```

```
In [49]: model3.score(X_test,y_test)
```

```
Out[49]: 0.7727272727272727
```

```
In [50]: #Applying Random Forest
from sklearn.ensemble import RandomForestClassifier
model4 = RandomForestClassifier(n_estimators=11)
model4.fit(X_train,y_train)
```

```
Out[50]: ▼ RandomForestClassifier
RandomForestClassifier(n_estimators=11)
```

```
In [51]: model4.score(X_train,y_train)
```

```
Out[51]: 0.993485342019544
```

```
In [52]: model4.score(X_test,y_test)
```

```
Out[52]: 0.7207792207792207
```

```
In [53]: #Support Vector Classifier
```

```
from sklearn.svm import SVC
model5 = SVC(kernel='rbf',
              gamma='auto')
model5.fit(X_train,y_train)
```

```
Out[53]: ▼ SVC
SVC(gamma='auto')
```

```
In [57]: model5.score(X_train,y_train)
```

Out[57]: 1.0

```
In [58]: model5.score(X_test,y_test)
```

Out[58]: 0.6168831168831169

```
In [59]: #Applying K-NN
from sklearn.neighbors import KNeighborsClassifier
model2 = KNeighborsClassifier(n_neighbors=7,
                             metric='minkowski',
                             p = 2)
model2.fit(X_train,y_train)
```

Out[59]: **KNeighborsClassifier**
KNeighborsClassifier(n_neighbors=7)


```
In [60]: #Preparing ROC Curve (Receiver Operating Characteristics Curve)
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(label, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
fpr, tpr, thresholds = roc_curve(label, probs)
print("True Positive Rate - {}, False Positive Rate - {} Thresholds - {}".format(tpr,fpr,thresholds))
# plot no skill
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
```

AUC: 0.836

```
True Positive Rate - [0.          0.06716418 0.23880597 0.44776119 0.60074627 0.75373134
0.88059701 0.98507463 1.          ], False Positive Rate - [0.          0.          0.02  0.056 0.12  0.248 0.428 0.668 1.
] Thresholds - [2.          1.          0.85714286 0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.          ]
```

Out[60]: Text(0, 0.5, 'True Positive Rate')



```
In [61]: #Precision Recall Curve for Logistic Regression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
```




```

yhat = model.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.624 auc=0.726 ap=0.727

Out[61]: [



In [62]: #Precision Recall Curve for KNN


```

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model2.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model2.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.658 auc=0.752 ap=0.709

Out[62]: [




```
In [63]: #Precision Recall Curve for Decision Tree Classifier

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model3.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model3.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='.')

```

f1=0.714 auc=0.815 ap=0.768

Out[63]: [




```
In [64]: #Precision Recall Curve for Random Forest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs = model4.predict_proba(features)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# predict class values
yhat = model4.predict(features)
# calculate precision-recall curve
precision, recall, thresholds = precision_recall_curve(label, probs)
# calculate F1 score
f1 = f1_score(label, yhat)
# calculate precision-recall AUC
auc = auc(recall, precision)
# calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3f auc=%.3f ap=%.3f' % (f1, auc, ap))
# plot no skill
plt.plot([0, 1], [0.5, 0.5], linestyle='--')
# plot the precision-recall curve for the model
plt.plot(recall, precision, marker='o')
```

f1=0.908 auc=0.963 ap=0.955

Out[64]: [



Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js